

I'll help you prepare for your interview by providing comprehensive questions and answers based on your research paper. Let me organize these by topic area.

1. FUNDAMENTAL CONCEPTS & MOTIVATION

Q: Can you explain what dimensionality reduction is and why it's important for metagenomic data?

A: Dimensionality reduction transforms high-dimensional data into lower-dimensional representations while preserving important information. In metagenomics, we deal with thousands of microbial features but only dozens to hundreds of samples, creating a "curse of dimensionality" problem. DR is essential because it: (1) reduces computational complexity, (2) enables visualization in 2D/3D, (3) reduces noise and prevents overfitting, (4) reveals underlying biological patterns like disease states or environmental gradients, and (5) facilitates downstream analyses like clustering and classification.

Q: What are the three main challenges in metagenomic data that you addressed?

A: The three key challenges are:

1. **High dimensionality:** Thousands of microbial features relative to sample size
2. **Extreme sparsity:** 90-99% zero entries due to zero-inflation from multiple biological and technical processes
3. **Compositionality:** Relative abundances sum to a constant (100%), creating mathematical dependencies between features that violate standard statistical assumptions

Q: Why is compositionality a problem for standard statistical methods?

A: Compositional data exists in a simplex space where values are constrained to sum to a constant. This creates spurious negative correlations—if one component increases, others must decrease to maintain the sum. Standard methods like Euclidean distance and correlation assume data in unrestricted Euclidean space, leading to misleading results. For example, two communities with identical absolute abundances but different total counts would appear artificially different.

2. PREPROCESSING METHODOLOGY

Q: Walk me through your preprocessing pipeline and justify each step.

A: Our three-step pipeline addresses both compositionality and sparsity:

Step 1 - nzCLR Transformation: We use non-zero Center Log-Ratio transformation that computes geometric means using only non-zero values and preserves zeros as NaN rather than using pseudocounts. This avoids arbitrary bias from pseudocount addition.

Step 2 - Iterative Matrix Completion: We impute the NaN values using scikit-learn's IterativeImputer (5 iterations), which models each feature as a function of others, capturing complex dependencies in microbiome structure.

Step 3 - Re-centering: Matrix completion disrupts the CLR constraint (sum of CLR values = 0). We restore this by subtracting row-wise means, ensuring the data maintains theoretical CLR properties.

This approach is more robust than simple pseudocount addition or multiplicative replacement.

Q: Why did you choose iterative imputation over other methods?

A: Iterative imputation (IterativeImputer) has several advantages:

1. It models complex dependencies between features rather than treating them independently
2. It doesn't introduce arbitrary bias like fixed pseudocounts
3. It's based on the MICE (Multivariate Imputation by Chained Equations) algorithm, which has theoretical grounding
4. It performed well in recent microbiome studies like DEICODE
5. After imputation, we re-center to preserve the compositional constraint

Q: What's the difference between standard CLR and nzCLR?

A: Standard CLR computes: $CLR(x_j) = \log(x_j / \text{geometric_mean}(\text{all values}))$. This fails when zeros exist because $\log(0)$ is undefined.

nzCLR modifies this: it computes the geometric mean using *only non-zero values* and marks zeros as NaN instead of forcing a transformation. This preserves the information that certain taxa are truly absent while avoiding the mathematical impossibility of $\log(0)$.

3. DIMENSIONALITY REDUCTION METHODS

Q: Why did you select these specific 8 DR methods?

A: We selected methods spanning diverse paradigms:

- **Linear methods** (PCA, PCoA): Computationally efficient baselines
- **Distance-preserving** (MDS): Global distance preservation
- **Local structure** (t-SNE): Excellent for discrete clusters
- **Topological** (UMAP): Balances local and global structure
- **Continuous structures** (PHATE): Preserves developmental/transitional patterns
- **Balanced approaches** (PaCMAP, SONG): Handle mixed discrete-continuous topologies

This diversity allows comprehensive comparison across different biological scenarios.

Q: What's the key difference between t-SNE and UMAP?

A:

- **t-SNE:** Focuses on local neighborhoods, uses probability distributions (Gaussian in high-D, t-distribution in low-D), computationally expensive ($O(n^2)$), primarily preserves local structure, can fragment global relationships
- **UMAP:** Uses topological data analysis (Riemannian geometry), more computationally efficient, better preserves global structure, has theoretical foundations in manifold theory, faster and scales better to large datasets

Q: When would you choose PHATE over other methods?

A: PHATE is ideal when your biological question involves:

1. **Continuous trajectories:** Disease progression, ecological succession, development
2. **Branching structures:** Cell differentiation, evolutionary divergence
3. **Gradients:** Environmental transitions (pH, temperature, depth)

PHATE uses diffusion processes that reveal smooth progressions better than t-SNE (which creates discrete clusters) or PCA (which may miss nonlinear patterns). In our Tara Oceans data, PHATE effectively revealed temperature gradients.

Q: Why did you use Jaccard distances specifically for PCoA?

A: Jaccard distance works on presence/absence data (binary), making it:

1. Robust to differences in absolute abundance or sequencing depth
2. Appropriate for phylogenetic/taxonomic similarity
3. Not affected by compositionality (it's scale-invariant)
4. Well-established in ecology for beta-diversity

We applied PCoA to raw counts converted to Jaccard distances rather than CLR-transformed data because Jaccard is specifically designed for binary ecological data.

4. EVALUATION METRICS

Q: Explain the difference between trustworthiness and continuity.

A: Both measure different aspects of neighborhood preservation:

Trustworthiness: "Are my neighbors in the low-D space trustworthy?" It penalizes points that appear as neighbors in the embedding but were NOT neighbors in high-D space (false neighbors). High trustworthiness means you can trust the clusters you see.

Continuity: "Did I lose my original neighbors?" It penalizes points that were neighbors in high-D but are NOT neighbors in low-D (missing neighbors). High continuity means the embedding didn't tear apart related samples.

Ideally, both should be high, but there's often a trade-off.

Q: Why did you include multiple distance metrics (Bray-Curtis, Jaccard, Aitchison)?

A: Each captures different biological aspects:

1. **Bray-Curtis:** Abundance-based dissimilarity, sensitive to dominant taxa, standard in ecology
2. **Jaccard:** Presence/absence, emphasizes species richness/diversity, robust to abundance differences
3. **Aitchison:** The mathematically correct distance for compositional data, respects simplex geometry

Different research questions require different metrics—studying keystone species needs abundance (Bray-Curtis), while species turnover needs presence/absence (Jaccard).

Q: What does normalized stress measure and why is lower better?

A: Normalized stress measures how well pairwise distances are preserved between high-D and low-D spaces. It's the ratio of distance discrepancies to total high-D distances. Lower values (closer to 0) mean distances are better preserved—samples that were far apart remain far apart, and close samples remain close. It's particularly important for methods like MDS that explicitly aim to preserve distances.

5. RESULTS & INTERPRETATION

Q: What was your main finding?

A: **No single DR method is universally optimal.** Method effectiveness depends on:

1. The specific research question (discrete groups vs. continuous gradients)
2. Dataset characteristics (sparsity, sample size, ecological context)
3. What needs to be preserved (local neighborhoods, global topology, or specific biological distances)

For example, t-SNE excelled at separating disease states in gut microbiome but fragmented continuous environmental gradients in ocean data, while PHATE did the opposite.

Q: Based on Table I, which methods performed best overall?

A: Different methods excel at different criteria:

- **Local structure:** t-SNE (all three local metrics), SONG and PaCMAP (trustworthiness)
- **Global structure:** MDS (normalized stress)
- **Ecological distances:** Jaccard PCoA (Bray-Curtis and Jaccard correlation)
- **Compositional structure:** PCA (Aitchison correlation)

This reinforces that method selection should be question-driven, not based on a single "best" method.

Q: Why did SONG perform well in your study?

A: SONG balances multiple objectives:

1. Preserves discrete clusters (like t-SNE/UMAP)
2. Maintains continuous progressions (like PHATE)
3. Shows connectivity between related clusters
4. Reveals branching structures without forcing artificial continuity
5. Achieved high trustworthiness (≥ 0.90 threshold)

This makes it particularly robust for exploratory analysis when you don't know if your data contains discrete groups, gradients, or both—which is common in real metagenomic studies.

6. TECHNICAL DETAILS

Q: Why did you set $k=10$ for neighborhood-based metrics?

A: $k=10$ is a standard choice that balances:

1. **Stability:** Large enough to be robust to noise
2. **Locality:** Small enough to capture local structure
3. **Literature precedent:** Commonly used in DR evaluation papers
4. **Biological relevance:** Captures meaningful local microbial community similarity

We could perform sensitivity analysis across different k values, but $k=10$ provides a reasonable default.

Q: What were your computational considerations?

A: We used:

- **Python 3.x** ecosystem for reproducibility
- **Parallel processing** where available (MDS, PHATE with `n_jobs=-1`)
- **Fixed random seeds** (`random_state=42`) for reproducibility
- **Departmental server infrastructure** for computational resources

Key packages: scikit-learn, umap-learn, pacmap, phate, song-vis, scikit-bio

Q: How did you handle the different input requirements for different methods?

A: We tailored inputs based on each method's assumptions:

- **Most methods** (PCA, MDS, t-SNE, UMAP, PaCMAP, PHATE, SONG): Used CLR-transformed data (Euclidean-based)
- **Jaccard PCoA**: Used raw counts converted to binary presence/absence, then Jaccard distance matrix

This ensures each method receives mathematically appropriate input rather than forcing all methods to use the same preprocessing.

7. DATASETS

Q: Why did you choose these three specific datasets?

A: They represent diverse ecological contexts:

1. **Human Gut (3,513 taxa)**: Host-associated, disease-related, high diversity
2. **PotatoSCMP Soil (1,620 taxa)**: Environmental gradient (pH), agricultural relevance
3. **Tara Oceans Marine (1,441 taxa)**: Global scale, temperature/depth gradients, medium diversity

This diversity tests DR method robustness across different environmental conditions, sample sizes, and biological questions.

Q: How did the datasets differ in terms of structure?

A:

- **Human Gut**: Clear discrete disease states vs. healthy controls—ideal for testing cluster separation (t-SNE, UMAP)
- **PotatoSCMP**: Continuous pH gradient—tests methods' ability to preserve environmental continuums (PHATE, PaCMAP)
- **Tara Oceans**: Temperature gradients across ocean basins—tests global geographic structure and environmental transitions

8. LIMITATIONS & FUTURE WORK

Q: What are the limitations of your study?

A: Several limitations exist:

1. **Parameter sensitivity**: We used default parameters; comprehensive hyperparameter tuning might reveal different rankings
2. **Limited datasets**: Three datasets may not cover all microbiome contexts (e.g., time-series, multi-omics)

3. **2D focus:** We only evaluated 2D embeddings; higher dimensions might perform differently
4. **Evaluation metrics:** No single ground truth exists; our metrics may not capture all biologically relevant aspects
5. **Computational cost:** We didn't systematically compare runtime/scalability

Q: What future work would you recommend?

A: Several promising directions:

1. **Extended evaluation:** More datasets (time-series, multi-kingdom, host-microbe interactions)
 2. **Hyperparameter optimization:** Systematic tuning for each dataset type
 3. **Integration methods:** Combining multiple DR methods or multi-view approaches
 4. **Deep learning:** Autoencoders, VAEs for compositional data
 5. **Downstream impact:** How DR choice affects clustering, classification, differential abundance
 6. **Interactive visualization:** Tools for biologists to explore multiple DR results simultaneously
 7. **Computational benchmarking:** Formal runtime and scalability analysis
-

9. PRACTICAL APPLICATIONS

Q: Based on your findings, what DR method would you recommend for a researcher studying IBD vs. healthy gut microbiome?

A: For discrete disease states (IBD vs. healthy):

1. **First choice: t-SNE with CLR transformation**—Excellent cluster separation, high trustworthiness and continuity
2. **Alternative: SONG or PaCMAP**—If you want to preserve potential disease progression gradients
3. **Complementary: PCA**—For interpretability and identifying major variance drivers

I'd recommend visualizing with multiple methods to ensure findings aren't method-specific artifacts.

Q: What about studying soil microbiome changes along an elevation gradient?

A: For continuous environmental gradients:

1. **First choice: PHATE**—Designed for continuous trajectories
2. **Alternative: PaCMAP**—Balances local clusters with global gradient structure

3. **Also consider: SONG**—If there might be discrete soil types along the gradient

Use correlation with Aitchison distance to verify the gradient is preserved.

10. BROADER CONTEXT

Q: How does your work contribute to the field?

A: Our contributions include:

1. **Comprehensive framework:** First systematic comparison of modern DR methods (including SONG, PaCMAP) specifically for metagenomics
2. **Robust preprocessing:** Integration of nzCLR with iterative imputation addresses both compositionality and sparsity
3. **Multi-metric evaluation:** Goes beyond simple visual inspection to quantitative validation
4. **Practical guidance:** Helps researchers select appropriate methods for specific biological questions
5. **Dataset diversity:** Tests across human, soil, and marine microbiomes

Q: How does this relate to recent work like DEICODE?

A: DEICODE (Robust Aitchison PCA) addresses compositionality and sparsity through matrix completion, similar to our approach. However:

1. **DEICODE focuses on PCA;** we compare 8 diverse methods
2. **We use nzCLR + iterative imputation;** DEICODE uses robust PCA with matrix completion
3. **We evaluate modern methods** (SONG, PaCMAP) not covered in DEICODE literature
4. **Our framework is method-agnostic;** it could incorporate DEICODE as preprocessing

Our work complements DEICODE by showing when alternatives to PCA might be preferable.

Q: What impact could this have on microbiome research?

A: Practical impacts include:

1. **Better visualization:** Researchers can choose methods matching their biological questions
 2. **Avoided artifacts:** Understanding trade-offs prevents misinterpretation
 3. **Standardization:** Evaluation framework could become standard practice
 4. **Tool development:** Guides creation of microbiome-specific visualization tools
 5. **Hypothesis generation:** Better visualizations lead to better biological insights
-

11. CHALLENGING/CRITICAL QUESTIONS

Q: Your preprocessing involves three steps that each transform the data. How do you know you're not introducing artifacts?

A: Valid concern. We address this through:

1. **Mathematical grounding:** Each step preserves theoretical properties (CLR constraint, compositional geometry)
2. **Comparative evaluation:** We compare against PCoA on raw data (Jaccard), which showed different but complementary patterns
3. **Biological validation:** Metadata correlations (disease, pH, temperature) align with known biology
4. **Literature support:** nzCLR and matrix completion are established in recent microbiome studies (DEICODE)
5. **Re-centering step:** Explicitly restores the CLR constraint

However, validation with mock communities or simulated data would strengthen this claim.

Q: Table I shows no single method meets all thresholds. Doesn't this make your findings less useful?

A: Actually, this is the **key finding**—it shows the importance of question-driven method selection rather than blind application of a single "best" method. Researchers should:

1. **Define objectives:** What matters for your question—local clusters, global structure, or specific distances?
2. **Select accordingly:** t-SNE for discrete groups, PHATE for gradients, PCoA for phylogenetic relationships
3. **Use multiple views:** Visualize with 2-3 methods to ensure robustness

This nuanced guidance is more valuable than falsely claiming one method is always best.

Q: You acknowledge using Claude AI for grammar. Did AI contribute to the scientific content?

A: No, AI was used only for grammar and language polishing, not for:

- Experimental design
- Data analysis
- Result interpretation
- Scientific conclusions
- Novel contributions

All scientific work, methodology, code implementation, and intellectual contributions are entirely our own. This is transparent acknowledgment of writing assistance, which is increasingly standard practice.

12. FUTURE DIRECTIONS QUESTIONS

Q: If you could extend this work, what would be your top priority?

A: **Downstream analysis impact study:** Quantify how DR method choice affects:

1. Clustering accuracy (compared to metadata labels)
2. Differential abundance detection
3. Classification performance
4. Biological interpretation consistency

This would provide actionable guidance: "Use method X if your downstream task is Y."

Q: How would you adapt this framework for time-series metagenomic data?

A: Time-series requires:

1. **Temporal metrics:** Evaluate trajectory smoothness, temporal coherence
2. **Dynamic DR methods:** Consider methods that explicitly model time (diffusion maps, PHATE with temporal weighting)
3. **Longitudinal structure:** Metrics for preserving subject-specific trajectories
4. **Evaluation adaptation:** Within-subject vs. between-subject distance preservation

This would be valuable for studies of disease progression or ecosystem dynamics.

Would you like me to:

1. Elaborate on any specific question?
2. Add more questions on a particular topic?
3. Provide practice answers for you to rehearse?
4. Create a summary of key talking points?